

ARIEL A. ARÉVALO ALVARADO

Founding AI Architect

arielarevaloalvarado@gmail.com ◇ (+506) 7064 1684 ◇ Heredia, Costa Rica ◇ LinkedIn ◇ GitHub

“Built 2 enterprise AI platforms from scratch as founding/primary engineer, delivering 50+ production agents and a product concept validated by a \$9M investment from OpenAI in a similar initiative. Engaged in parallel practice in empirical AI evaluation and AI safety research training.”

WORK HISTORY

Chief AI Architect (Founding) · CRSS, Remote

Jun 2025 – Present

- Established an AI engineering organization spanning 33 repositories and 10 engineers, including production platforms, client prototypes, documentation hubs, and team training resources.
- Designed an apprentice agent workflow platform (AutoActivity) from zero to production as the chief architect, delivering a 4-repository system spanning backend (FastAPI), frontend (Playwright browser event trace capture), and Azure ML data pipelines.
- Architected and shipped the TracePipe pivot, extracting AutoActivity’s browser-trace capture, processing, and canonical-action classification into a standalone product after a competitor raised \$9M from OpenAI for the same concept.
- Architected and built across 5+ client-specific AI prototypes (chatbots, document classifiers, and data analysis agents) using LangChain, MCP, Azure AI Foundry, and Slack integrations.

Skills: Python, FastAPI, Azure ML, LangChain, Docker, MCP, Alembic, Azure AI Foundry, System Architecture

Senior AI Engineer · Colibri Group, Remote

Jul 2025 – Present

- Initiated the Colibri AI Foundations enterprise strategy, autonomously building an in-house stack across infrastructure (model routing, key serving, telemetry), platform (*Instant Agents*), and tooling (harness, observability). Company leadership promoted it to a dedicated team and initiative, which I now lead.
- Designed a zero-deployment agent configuration architecture (*Instant Agents*) that reduced per-agent deployment effort by 45% through a control plane/data plane separation with immutable versioning and automated component discovery. Scaled from 6 to 50+ agents at \$8.52/agent/month with under 1% infrastructure overhead.
- Built the Rubi enterprise AI platform from scratch as the primary developer, implementing agent runtime orchestration, JWT security, persistent memory, and a self-serve agent builder on AWS Bedrock AgentCore.

Skills: Python, AWS Bedrock AgentCore, MCP, Abacus AI, System Architecture

AI/ML Engineer · Datasite, Heredia

Apr 2024 – Jun 2025

- Designed the first end-to-end MLOps pipeline at Datasite, replacing ad-hoc Jupyter notebook experimentation with reproducible Azure ML pipeline steps incorporating data versioning and experiment tracking.
- Architected a medallion data architecture (Bronze/Silver/Gold) for production ML workflows, establishing standardized ingestion, validation, and feature engineering stages on Airflow and Azure ML.
- Maintained and extended existing production systems running classical NLP methods (NER, PII detection, BM25 search), regression models with Scikit-learn, and deep learning classifiers (MLPs) using PyTorch and Keras.

Skills: Python, Pandas, Azure ML, Apache Airflow, MLOps, Scikit-learn, Keras, PyTorch, NLP, Deep Learning

Data Engineer · TransUnion, Heredia

Dec 2022 – Apr 2024

- Led development of the initial ETL pipeline synthesizing IPv6 geolocation data for the industry-leading IP GeoPoint product, designing Spark SQL and Spark ML workflows on Google Cloud Dataproc.
- Maintained an ETL pipeline synthesizing IP reputation data to determine whether users were malicious and/or bots, making use of a Decision Tree model running in J2EE.

Skills: Scala, Java, Python, Apache Spark, Spring Framework, GCP Dataproc

Software Developer · Snap Finance, Alajuela

Dec 2021 – Dec 2022

- Developed originations back-end services across monolith (J2EE), microservices (Spring), and serverless (AWS Lambda) architectures for consumer financing applications.

Skills: Java, JavaScript, Python, Spring Framework, J2EE, AWS Lambda

PROJECTS

Moral Licensing Comprehension[↗]: Measuring Flexibility of Social Prediction in LLMs

Apr 2026

- Solo submission to the Google DeepMind/Kaggle *Measuring Progress Toward AGI* competition (Social Cognition track). Tested whether LLMs can reason about moral licensing: do they predict that past virtue *frees* rather than *binds* future behavior in third-party agents? Used a 2×2 within-subjects factorial (120 items, 60 matched pairs), with ground truth grounded in the Blanken et al. (2015) meta-analysis.
- Built the autonomous research harness that produced this submission alongside 13 other hypothesis-driven benchmarks, with formal statistics (Cochran's Q, Holm-corrected pairwise, within-tier discrimination) over a 5-day deadline.

Skills: Python, Claude Agent SDK, OpenRouter, Pytest, Statistical Analysis, Empirical Evaluation

Rigger[↗]: Declarative coding-agent harness framework

Mar 2026

- Built and released as Apache-2.0 OSS on PyPI a Python framework for declarative coding-agent harness management, with users defining harnesses as code or configuration and the framework managing their development lifecycle from workflow specifications.

Skills: Python, Claude Agent SDK, Automated Software Engineering, Evolutionary Algorithms

ClimateFact[↗]: AI-powered climate statement fact-checker

Spring 2025

- Built an AI fact-checking system using LangGraph and Azure OpenAI to detect contradictions in climate statements against the IPCC AR6 report, leveraging DeBERTa natural language inference models for semantic validation.

Skills: Python, LangGraph, Azure OpenAI, NLP, NLI, DeBERTa

PUBLICATIONS

- K. Villalobos-Solís, E. Viquez-Mora, A. Badilla-Olivas, E. Vilchez-Lizano, and **A. Arévalo-Alvarado (Senior Author)**, “Multimodal Foundation Models for Zero-Shot User Activity Classification”, *JOCICI 2026* [↗].

PROFESSIONAL DEVELOPMENT

BlueDot Impact, AGI Strategy

Feb – Mar 2026

BlueDot Impact, Technical AI Safety

Apr – May 2026

AAAI Workshop, Towards Execution-Grounded Automated AI Research

Feb 2026